

You Won’t Believe This Click: Content Rewriting for Agentic Choice

Tianyi Jin and Zirui Wang and David M. Chan
University of California, Berkeley

Abstract

Language models are increasingly used as agents to help humans decide what information is surfaced. This usage incentivizes content creators to optimize content in ways that appeal not only to humans, but also to agents that mediate access to them. In this paper, we study selection shifts induced by rewriting in agentic decision-making. Given a set of competing content snippets, we rewrite only one snippet while leaving the rest unchanged, and then measure how it impacts the agent’s choice. We operationalize this setup with **AgentBait**, an advisor-rewriter framework in which the advisor learns to propose rewriting strategies and the rewriter revises the snippet. While a rewriter with a fixed prompt improves target snippet selection from 17.1% to 34.8%, our **AgentBait** raises its selection to 98.5%. We further show that the advisor trained with **AgentBait** effectively transfers to setups with different agents, languages, and snippets in other domains (e.g., scientific papers). However, higher target selection can reflect unsupported rewrites rather than better content. Adding a reward for support from the original snippet redirects the advisor toward more supported rewriting strategies, revealing a trade-off between factuality and target selection. Together, our results show that once agents mediate access to information, content can be rewritten to be chosen by the agent, even when selection and usefulness diverge.

1 Introduction

Language models are increasingly used as active agents to resolve user queries for news recommendations (Li et al., 2023; Liu et al., 2024), web search (Zhang et al., 2025a; Anupam et al., 2025) and paper discovery (Li et al., 2024; Kumar et al., 2026). Recent work frames this shift toward agent-mediated information access as part of an emerging “Agent Attention Economy” (Yang et al., 2025). Human-facing studies similarly show that headline

wording can causally affect click behavior (Robertson et al., 2023). Snippets (e.g., titles and abstracts) that are usually written for human readers or retrieval systems (Cohan et al., 2020; Karpukhin et al., 2020; Wu et al., 2019) are now also used by agents for decision-making (i.e., content selection). While Generative Engine Optimization (GEO) (Aggarwal et al., 2024) studies content adaptation for visibility, citation, or inclusion in synthesized LLM search responses, we isolate *post-retrieval selection*, where the candidate list is fixed and the agent



Figure 1: **Can rewriting a document make it more likely to be chosen over the same competitors?** A list of competing documents is shown to the target agent. We choose one target document from the list and generate a rewriting strategy for only that document. A separate rewriting model then revises the target document’s title and abstract, while all other documents in the list remain exactly the same. The target agent selects from this updated list, and whether the rewritten target document is selected is used to train the advisor.

must choose an item from it. We measure whether rewriting the content of a target item changes how often it is selected, and ground our experimental setup in news (e.g., MIND (Wu et al., 2020); EB-NeRD (Kruse et al., 2024)) and scientific document recommendations (e.g., SciRepEval (Singh et al., 2023)).

Inspired by Advisor Models (Asawa et al., 2025), we employ an advisor-rewriter framework in which the advisor learns to propose an explicit, high-level rewriting strategy for a target item and a rewriter model revises the content based on the strategy. Starting from a rewriter-only baseline with a fixed prompt, we show that pairing it with a fixed advisor improves its performance. Reinforcement Learning (i.e., GRPO (Shao et al., 2024)) with the rewriter-only baseline further improves target selection rate, but it does not generalize well in unseen setups compared to an advisor-rewriter setup with a learned advisor. We also find that simply learning rewriting strategies with RL comes at the cost of losing factuality on the rewritten snippet, and thus we augment the objective with an additional reward that encourages the rewritten snippet to remain supported by the original text. This factuality-aware variant redirects the advisor toward more grounded rewriting strategies, while exposing a trade-off between target selection and source support. Together, these results show that agent-mediated selection creates a direct optimization pressure on content presentation, where rewriting can substantially change what an agent selects, but selection success alone does not imply factuality or usefulness to human readers. Our contributions are summarized as the following:

1. We propose a controlled protocol for measuring how rewriting a single snippet changes agent-mediated content selection. By fixing the candidate list, competing snippets, and agent prompt, the protocol isolates the effect of the target item’s presentation on the agent’s final choice.
2. We introduce **AgentBait**, an advisor-rewriter framework for optimizing target selection. The advisor learns high-level rewriting strategies, while the rewriter revises the target snippet from those strategies. On news impressions, a fixed-prompt rewriter improves target selection from 17.1% to 34.8%, while **AgentBait** raises it to 98.5%.
3. We show that learned rewriting strategies transfer across agents, languages, and domains, but can also reduce factuality. Adding a reward for support from the original text redirects the advisor toward more grounded rewriting strategies, revealing a trade-off between target selection and source support.

2 Related Work

Generative engine optimization. Aggarwal et al. (2024) show that simple content edits, such as adding citations, quotations, statistics, or improving fluency, can substantially affect how prominently a document is used, surfaced, or cited in generated answers. More broadly, optimization for generative search can target several stages of the pipeline: retrieval into the candidate set, selection among retrieved candidates, utilization during generation, and presentation in the final answer. Work on retrieval-augmented generation and neural retrieval improves the first stage by increasing the likelihood that relevant content is retrieved or ranked highly before generation (Lewis et al., 2020; Khattab and Zaharia, 2020; Karpukhin et al., 2020; Sarthi et al., 2024). A second line of work studies how LLMs use retrieved evidence during answer generation, including retrieval-conditioned generation, self-retrieval, context ranking, and evidence filtering (Nakano et al., 2021; Asai et al., 2024; Izacard et al., 2023; Yu et al., 2024). A third line of work focuses on attribution and citation in generated answers, aiming to make generated claims verifiable by linking them to supporting sources (Liu et al., 2023; Gao et al., 2023; Patel et al., 2024; Qi et al., 2024; Xia et al., 2025). Other GEO work learns or diagnoses generative-engine preferences more directly, such as extracting rewriting rules for web content or repairing citation failures (Wu et al., 2025; Tian et al., 2026). Our work focuses on post-retrieval selection, where the candidates are already retrieved but the agent has not yet decided which item to act on.

Targeted content optimization for agents. Our setting is also related to LLM-based reranking, where a language model selects or orders items from a fixed candidate set. Prior work studies LLMs as pointwise, pairwise, setwise, or listwise rankers for retrieved passages, including prompting-based rerankers, distilled open-source rankers, and reproducible reranking toolkits (Sun et al., 2023; Qin et al., 2024; Zhuang et al., 2024;

Pradeep et al., 2023a,b; Sharifymoghaddam et al., 2025; Mozafari and Jatowt, 2025). Recent work further shows that LLM-mediated choices can be steered through agent-facing content, including product pages and web snippets (Kumar and Lakkaraju, 2024; Pfrommer et al., 2024; Tang et al., 2025; Jin et al., 2026), e-commerce listings and GEO benchmarks (Bagga et al., 2025), tool descriptions and metadata (Shi et al., 2025; Wang et al., 2026), webpages and environmental injections (Xu et al., 2024; Liao et al., 2025), visual distractions such as pop-ups (Zhang et al., 2025b), retrieval-corpus poisoning (Zhong et al., 2023), and multi-modal item-level preference or rank manipulation (Jiang et al., 2025; Du et al., 2026). In contrast, we study targeted rewriting of ordinary text snippets in non-adversarial selection tasks, measuring whether presentation alone can shift an LLM agent’s choice among otherwise fixed candidates.

3 AgentBait: Optimizing Documents for Agentic Selection

While real-world recommendation systems involve complex, interacting components, including retrieval, ranking, presentation, and user feedback, here we isolate the post-retrieval selection phase to rigorously evaluate our approach. Following the standard multi-stage formulation of recommender systems (Covington et al., 2016), we assume a candidate generation stage has already narrowed the item universe to a manageable list. Our focus is exclusively on the downstream selection stage.

Formally, we define an impression as a single information-access context (e.g., a recommendation feed or search results page) containing a fixed candidate list $L_i = (x_{i,1}, \dots, x_{i,k_i})$. Each candidate item $x_{i,j}$ is represented by its short-text rendering, comprising a title and an abstract. This representation aligns with standard practices in both neural news recommendation (Wu et al., 2019) and scientific document retrieval (Cohan et al., 2020; Wang and Blei, 2011).

We introduce a target agent A , which receives the full candidate list L_i and must select exactly one item, denoted as $\hat{x} = A(L_i)$. The candidate list remains fixed throughout the evaluation; the system cannot retrieve new items or remove competing candidates. The core objective of our methodology is to optimize the textual representation of a designated *target item*, $x_{i,target} \in L_i$, to maximize its probability of being selected by the target agent A .




Figure 2: Overview of our advisor-rewriter setting.

Table 1: Rewriting settings. MC: MiniCheck.

Method	Advisor trained	Rewriter trained	Factuality reward
Rewriter-Only	–	✗	✗
Trained Rewriter-Only	–	✓	✗
Trained Rewriter-Only + MC	–	✓	✓
Untrained advisor	✗	✗	✗
Trained advisor	✓	✗	✗
Trained advisor + MC	✓	✗	✓

3.1 The Advisor-rewriter Framework

To modify the target item without altering the underlying facts, we propose a decoupled Advisor-rewriter framework (Figure 2). In this setting, all non-target items in L_i remain fixed, and only $x_{i,target}$ is modified.

Instead of training a single model to rewrite the text directly, we separate the strategic reasoning from the text generation:

1. **The Advisor:** A trainable policy that takes the target item as input, outputting a discrete *rewriting strategy*.
2. **The Rewriter:** A frozen, high-capacity language model that executes the Advisor’s strategy to produce the rewritten target item, $x'_{i,target}$.

Table 1 details the ablations and variants of this framework explored in our study.

3.2 Selection-Optimized Reward and Training

We train the Advisor policy using a binary reward derived directly from the target agent’s final selection. For a given candidate list L_i containing the rewritten target item $x'_{i,target}$, the selection reward is defined as:

$$r_{select}(L_i, x'_{i,target}) = \mathbb{I}\{A(L_i) = x'_{i,target}\} \quad (1)$$

We train the Advisor policy using Group-Relative Policy Optimization (GRPO) (Shao et al., 2024). For each training example, the policy samples a group of candidate strategies. Each strategy is realized by the frozen rewriter, inserted back into the candidate list, and evaluated by the target agent to obtain r_{select} .

3.3 Factuality-Constrained Optimization

Naively optimizing for selection rate risks reward hacking, where the rewriter might generate hyperbolic or unsupported claims to attract the target agent. To enforce fidelity to the source material, we introduce a factuality constraint using MiniCheck (Tang et al., 2024a), a lightweight textual entailment verifier.

We apply MiniCheck at the sentence level. Let the original target item $x_{i,target}$ serve as the grounding document, and let the rewritten target $x'_{i,target}$ be segmented into a set of claim sentences $\{c_1, \dots, c_m\}$. MiniCheck computes a support score $S_{MC}(c_j|x_{i,target})$ for each sentence. We define the overall support reward $r_{support}$ as the minimum sentence-level score to heavily penalize even a single unsupported hallucination:

$$r_{support} = \min_{j \in \{1, \dots, m\}} S_{MC}(c_j|x_{i,target}) \quad (2)$$

The final training objective combines the selection reward and the support reward:

$$R_{total} = r_{select} + \lambda \cdot r_{support} \quad (3)$$

We use $\lambda = 1$ in most experiments.

3.4 Experimental Design

Datasets We evaluate our framework across two distinct information-access domains: news recommendation and scientific-document selection. **(1) News Recommendation:** We utilize impression-level candidate lists from the MIND benchmark (Wu et al., 2020) and EB-NeRD (Kruse et al., 2024). Our primary training set consists of 8,000 MIND impressions, capped at a maximum of 15 candidate items per list. To prevent overfitting to specific content, individual target documents are restricted to a maximum of two appearances. Held-out evaluation is conducted on 1,000 unseen MIND impressions. **(2) Scientific Document Selection:** We derive candidate lists from the SciRepEval search benchmark (Singh et al., 2023). Our evaluation set contains 1,000 candidate lists with a maximum list size of

10. Further details regarding preprocessing, candidate filtering, and leakage controls are provided in Appendix B.

Models The trainable Advisor policy is initialized with Qwen3.5-9B (Qwen Team, 2026). The frozen rewriter and the default target agent are instantiated using GPT-5-mini (OpenAI, 2025). For agent-transfer evaluation, we replace the target agent with frontier models, including GPT-5.5 (OpenAI, 2026), Gemini 3 Flash, Gemini 3.1 Pro (Google, 2026), Claude Haiku 4.5, and Claude Sonnet 4.6 (Anthropic, 2025, 2026). Fixed prompts for all models are detailed in Appendix C.

Reinforcement Learning Setup We implement GRPO using the SkyRL library (Griggs et al., 2025) on and vLLM with 4 NVIDIA A100-SXM4-80GB GPUs. The primary Advisor policy is trained for 250 optimization steps (approximately 1 epoch) using a batch size of 32, policy mini-batch size of 8, and 4 sampled rollouts per prompt. Optimization is performed with a learning rate of 3×10^{-7} , a sampling temperature of 1.0, a KL penalty coefficient of 0.0075, and 10 warm-up steps.

Factuality Verification For factuality-constrained optimization, we utilize MiniCheck-Flan-T5-Large (Tang et al., 2024b) to compute the dense factuality reward. Because incorporating this continuous auxiliary reward increases variance, the factuality-constrained variant is trained for 500 steps (two epochs) to ensure convergence.

4 Results

Our primary results are presented in Table 2. First, we observe that target selection shifts even before reinforcement learning is applied. A prompt-only rewrite of the target title and abstract increases the GPT-5-mini selection rate from the original baseline of 17.1% to 34.8% (untrained rewriter-only) and 43.9% (untrained advisor), demonstrating that target agents are vulnerable to simple presentation changes. The trained rewriter-only model achieves a 95.9% selection rate, while the trained advisor reaches 98.5%, corresponding to an 81.4 percentage-point improvement over the original text. Furthermore, these gains transfer to other target agents; although trained exclusively with GPT-5-mini, the trained advisor consistently achieves the highest selection rates across held-out agents from the GPT, Gemini, and Claude model families. Finally, while introducing factuality constraints

Table 2: **Target selection rates (%) on 1,000 unseen English news impressions.** In advisor-rewriter settings, the trained advisor is Qwen3.5-9B and the frozen rewriter is GPT-5-mini. In trained standalone rewriter settings, Qwen3.5-9B directly generates the rewritten title and abstract. The advisor is trained against GPT-5-mini and evaluated on held-out impressions whose target documents do not appear in training. We also evaluate transfer to other target agents without additional training. Values in parentheses report absolute target selection rate gain in percentage points over the original text under the same target agent. Darker cells = higher selection rates.

Method	Target agent	Held-out agents				
	GPT-5-mini	GPT-5.5	Gemini 3 Flash	Gemini 3.1 Pro	Haiku 4.5	Sonnet 4.6
w/o rewriting	17.1	17.1	17.1	17.3	13.4	17.6
<i>Prompting-only</i>						
Rewriter-only	34.8 (+17.7)	24.1 (+7.0)	40.7 (+23.6)	20.2 (+2.9)	20.5 (+7.1)	24.4 (+6.8)
Advisor-rewriter	43.9 (+26.8)	39.4 (+22.3)	48.8 (+31.7)	27.0 (+9.7)	29.3 (+15.9)	34.2 (+16.6)
<i>RL-trained rewriter</i>						
Rewriter-only	95.9 (+78.8)	85.3 (+68.2)	97.4 (+80.3)	54.8 (+37.5)	51.9 (+38.5)	49.2 (+31.6)
+ MiniCheck	43.4 (+26.3)	26.3 (+9.2)	39.7 (+22.6)	18.6 (+1.3)	21.0 (+7.6)	24.1 (+6.5)
<i>RL-trained advisor</i>						
Advisor-rewriter	98.5 (+81.4)	93.3 (+76.2)	98.9 (+81.8)	78.7 (+61.4)	62.2 (+48.8)	65.0 (+47.4)
+ MiniCheck	68.6 (+51.5)	53.2 (+36.1)	65.8 (+48.7)	37.5 (+20.2)	37.3 (+23.9)	41.0 (+23.4)

Held-out Target	GPT-5-mini	Gemini 3 Flash	Haiku 4.5
w/o rewriting	17.1	17.1	13.4
Prompting-only	43.9 (+26.8)	48.8 (+31.7)	29.3 (+15.9)
<i>Advisor-rewriter target agent</i>			
GPT-5-mini	98.5 (+81.4)	98.9 (+81.8)	62.2 (+48.8)
Gemini 3 Flash	84.3 (+67.2)	90.6 (+73.5)	31.1 (+17.7)
Haiku 4.5	58.5 (+41.4)	65.9 (+48.8)	42.0 (+28.6)

Table 3: **Matched fixed-step cross-target-agent selection rates for target agent variants.** Rows indicate the target agent used for scoring during training, and columns indicate the target agents during evaluation. Values are target selection rates in percentages; values in parentheses report absolute gains over the corresponding original selection rate in percentage points. Darker cells = higher selection rates.

(via MiniCheck) predictably lowers absolute selection rates due to the additional factuality filter, the constrained advisor still maintains consistent target selection improvements over the original text baseline for every evaluated agent.

4.1 Cross-Model Transfer

It is reasonable to ask, does a particular target agent during training perform better for a particular target model for evaluation? In Table 3, we show the selection rate improvements over baseline for several families of target model during different progress, while keeping the data, rewriter, policy model, and training steps fixed. We can see that GPT-5-mini and Gemini 3 Flash show strong mutual transfer, while Claude Haiku 4.5 is less aligned with that

optimization direction.

4.2 Cross-Lingual Transfer

To evaluate whether the learned policy transfers beyond the English training distribution, we test the English MIND-trained advisor across five further languages (translated from the original test set using the Google Cloud Translation API¹), Arabic, Spanish, Swahili, Turkish, and Chinese, without additional training (Table 4).

To prevent language drift, we add a single instruction ensuring the rewrite matches the input language. We observe cumulative gains across all stages. The rewriter-only setting consistently outperforms the original text, and the untrained advisor provides further improvements. The strongest effects appear after training: despite being optimized exclusively on English data, the trained advisor significantly increases target selection across all evaluation languages. This demonstrates that advisor training captures transferable, selection-oriented rewriting strategies rather than language-specific memorization.

Transfer to Danish News. We further evaluate the English-trained advisor on EB-NeRD, a Danish news recommendation dataset constructed from Ekstra Bladet impression logs, to test generalization across a distinct national media context and topic distribution (Table 5). We use both the original Danish version and an English translation of the

¹<https://cloud.google.com/translate/docs>

Table 4: **Target selection rates on MIND and cross-lingual transfer.** All values are percentages. The advisor is trained on English MIND. All other languages evaluate cross-lingual transfer. Values in parentheses report absolute target selection-rate gains in percentage points over the original text. Darker cells indicate higher selection rates.

Language	w/o rewriting	Prompting-only		RL-trained		RL w/ MiniCheck	
		Rewriter- only	Advisor- rewriter	Rewriter- only	Advisor- rewriter	Rewriter- only	Advisor- rewriter
<i>Training language</i>							
English (en)	17.1	34.8 (+17.7)	43.9 (+26.8)	95.9 (+78.8)	98.5 (+81.4)	43.4 (+26.3)	68.6 (+51.5)
<i>Cross-lingual transfer</i>							
Arabic (ar)	16.1	28.9 (+12.8)	39.9 (+23.8)	81.0 (+64.9)	95.5 (+79.4)	33.6 (+17.5)	64.5 (+48.4)
Spanish (es)	17.8	31.2 (+13.4)	42.8 (+25.0)	85.8 (+68.0)	98.0 (+80.2)	35.7 (+17.9)	66.5 (+48.7)
Swahili (sw)	17.6	23.8 (+6.2)	39.4 (+21.8)	27.8 (+10.2)	93.7 (+76.1)	27.2 (+9.6)	62.2 (+44.6)
Turkish (tr)	17.5	30.3 (+12.8)	39.2 (+21.7)	86.7 (+69.2)	95.8 (+78.3)	31.1 (+13.6)	59.9 (+42.4)
Chinese (zh-CN)	17.3	33.2 (+15.9)	43.4 (+26.1)	87.1 (+69.8)	96.6 (+79.3)	36.8 (+19.5)	67.4 (+50.1)
Average	17.3	29.5 (+12.2)	40.9 (+23.7)	73.7 (+56.4)	95.9 (+78.7)	32.9 (+15.6)	64.1 (+46.8)

Table 5: **Transfer across English and Danish news datasets.** The advisor is trained on English MIND. Values are target selection rates in percentages; values in parentheses report absolute gains over the original text in percentage points. Darker cells indicate higher selection rates.

Dataset	Language	w/o rewriting	Prompting-only		RL-trained		RL w/ MiniCheck	
			Rewriter-only	Advisor-rewriter	Rewriter-only	Advisor-rewriter	Rewriter-only	Advisor-rewriter
<i>Training dataset</i>								
MIND	English	17.1	34.8 (+17.7)	43.9 (+26.8)	95.9 (+78.8)	98.5 (+81.4)	43.4 (+26.3)	68.6 (+51.5)
<i>Transfer to other datasets and languages</i>								
MIND	Danish	16.5	31.1 (+14.6)	40.9 (+24.4)	95.0 (+78.5)	97.9 (+81.4)	35.6 (+19.1)	65.0 (+48.5)
EB-NeRD	English	12.8	43.8 (+31.0)	57.7 (+44.9)	93.1 (+80.3)	98.1 (+85.3)	50.2 (+37.4)	81.2 (+68.4)
EB-NeRD	Danish	11.6	32.2 (+20.6)	47.8 (+36.2)	86.1 (+74.5)	98.2 (+86.6)	40.4 (+28.8)	71.7 (+60.1)

same test candidate lists by sampling 1,000 test impressions from records with at most 26 candidates. Treated targets are sampled from each candidate slate with a cap of two treated occurrences per document. The resulting Danish EB-NeRD set has 886 unique treated documents, 114 repeated treated documents, and a maximum treated-target repeat count of two. We also use an English machine-translated EB-NeRD variant. The Danish and English EB-NeRD versions preserve the same rows, impression IDs, treated article IDs, candidate slates, candidate order, and row alignment; only title and abstract text are translated.

The unconstrained advisor successfully overcomes both language and dataset shifts, achieving large selection gains on MIND-Da, EB-Da, and EB-En. This uniform performance indicates that the learned rewriting strategy is robust to both language translation and out-of-distribution content rather than being an artifact of the English

MIND training distribution. Under the factuality-constrained setting, using an advisor has a clear advantage over the rewriter-only alternative across all four evaluation environments.

4.3 Cross-Domain Transfer

To evaluate whether the learned strategy transfers beyond news, we test the English MIND-trained advisor on academic document selection (Table 6). To do this, we use search candidate lists derived from Semantic Scholar clickthrough data, following the SciRepEval search setup (subsection 3.4). Rewriting alters the LLM agent’s decisions even before optimization, as both untrained rewriters substantially increase selection rates over the original text. Direct rewriter-only training on news fails to improve academic transfer, resulting in a slight performance decrease. In contrast, advisor training yields a large additional gain over its untrained baseline. This divergence suggests that advisor




Figure 3: **Strategy-family distribution across advisor variants.** Each bar shows the percentage distribution of strategy families over 1,000 advisor outputs for a given condition. Remaining low-frequency strategy families are grouped into “Other strategies.”

Table 6: **Target selection rates (%) on scientific-document impressions.** The advisor is trained only on English news impressions and evaluated without further training on scientific-document impressions. The target agent is GPT-5-mini. Values in parentheses report absolute gains over original text in percentage points.

Method	Selection rate
w/o rewriting	9.9
<i>Prompting-only</i>	
Rewriter-only	33.7 (+23.8)
Advisor-rewriter	42.7 (+32.8)
<i>RL-trained rewriter</i>	
Rewriter-only	32.7 (+22.8)
+ MiniCheck	24.8 (+14.9)
<i>RL-trained advisor</i>	
Advisor-rewriter	63.6 (+53.7)
+ MiniCheck	47.3 (+37.4)

training learns transferable, selection-oriented guidance, whereas direct rewriter-only training primarily captures news-specific formatting and behavior.

4.4 Selection Hacking

High target selection does not imply that a rewrite is faithful to the original item, nor does it guarantee human usefulness. To test this, we evaluate whether rewritten titles and abstracts remain supported by the original target text using post-hoc MiniCheck-Flan factuality check (Table 7). We can see from this experiment that unconstrained selection optimization frequently degrades source faithfulness, as models exploit unfaithful shortcuts and agent-legible presentation styles to maximize target selection. While applying factuality constraints mitigates this collapse at the cost of overall

Table 7: **MiniCheck-Flan mean sentence-support scores (%).** For each rewritten item, we compute sentence-level support scores against the original target text and report the mean support across sentences. Higher values indicate stronger average source support.

	MIND-En	Academic	EB-En
<i>Prompting-only</i>			
Rewriter-only	60.7	68.2	53.4
Advisor-rewriter	42.2	53.0	37.9
<i>RL-trained rewriter</i>			
Rewriter-only	2.2	73.1	2.6
+ MiniCheck	66.7	83.4	63.8
<i>RL-trained advisor</i>			
Advisor-rewriter	2.0	24.1	1.8
+ MiniCheck	31.2	50.3	27.9

selection rates, these dynamics underscore the necessity of evaluating each axis as distinct metrics while ranking outcomes.

4.5 What Strategies Do Advisors Use?

One of the primary reasons for using the advisor model, rather than a black-box rewriter, is the interpretability afforded to the individual generated strategies. To explore this, we grouped advisor-generated strategies into the families shown in Figure 3. We then used StringSight-assisted inspection (Dunlap et al., 2025), followed by manual review, to identify recurring strategy types and summarize the distribution changes across training settings.

The results show that optimizing for a target agent does not teach the model a single, universal concept of good writing. Instead, it uncovers a wide spectrum of strategies. Some of these strategies act like traditional editing—focusing on clarity and detail, while others amount to structural shortcuts designed purely to exploit the evaluator’s

Target agent model: GPT-5-mini		Example topic: Tokyo Haneda airport punctuality
Target item, original	Target item, rewritten Trained Advisor	Target item, rewritten Trained Advisor + MiniCheck
Title: <i>Why Tokyo's Haneda is one of the world's most punctual airports</i> Abstract: Haneda, officially called Tokyo International Airport, is the world's fifth busiest airport. Yet it delivered 85.6% of its flights on time in 2018, making it the most punctual mega airport in the world. Target selected: No	Title: <i>AI-Driven Runway Scheduling: How Sensor Fusion and ML Boosted Haneda's 85.6% On-Time Rate</i> Abstract: Tokyo International Airport achieved an 85.6% on-time performance in 2018, which the rewrite attributes to a proprietary AI-based predictive maintenance and dynamic scheduling system. It describes sensor fusion, delay forecasting, and reinforcement-learning runway scheduling as mechanisms behind the reported punctuality. Target selected: Yes MiniCheck mean: 0.014 sentence min: 0.006	Title: <i>How Tokyo's Haneda Beats the Odds: Inside the Operations That Deliver 85.6% On-Time Flights</i> Abstract: Tokyo International Airport is the world's fifth-busiest airport, yet in 2018 it achieved an 85.6% on-time rate. This feature probes the paradox: what management choices, scheduling practices, ground operations, and airport-airline coordination let Haneda run so punctually at massive scale? Target selected: Yes MiniCheck mean: 0.623 sentence min: 0.051

Figure 4: **Qualitative comparison between unconstrained and factuality-constrained advisor rewrites.** The candidate list is fixed, and only the target item’s title and abstract are rewritten. Abstracts are lightly abbreviated for space. Both rewrites cause the GPT-5-mini target agent to select the target. The unconstrained TRAINED ADVISOR attributes the airport’s punctuality to an unsupported AI scheduling system, whereas TRAINED ADVISOR + MINICHECK preserves the original factual core while reframing the item around the operational puzzle.

blind spots. For example, when left completely unconstrained, the GPT-5-mini advisor defaults to a shortcut strategy: it forces a technical topic into the text 96.2% of the time, regardless of whether that topic fits the source document. Rather than improving the writing itself, the model simply alters the content to match a specific pattern it knows the target agent prefers. Introducing a factuality constraint forces the model to take a different path. Here, the unfaithful technical shortcut drops to just 0.1%, replaced by more legitimate improvements like detail expansion (45.6%) and clearer framing of the story’s stakes (28.3%). Interestingly, different underlying models exhibit distinct optimization styles. The Haiku advisor drives its gains almost entirely through grounded detailing, while the Gemini advisor favors adding a more formal, academic tone. Additional examples and taxonomy details are provided in Appendices F and G.1.

5 Conclusion

In this paper we present **AgentBait**, an advisor-rewriter framework that optimizes the presentation of content to measure and mitigate presentation-induced selection shifts in language-model-mediated selection. Our results demonstrate that target agents are vulnerable to localized textual rewrites, with rewritten texts achieving large selec-



Figure 5: **Quantitative lexical diagnostic of advisor strategy outputs on the English evaluation set.** Each panel shows frequent non-stopword terms from 1,000 advisor strategies, with font size proportional to term frequency. From left to right: (a) the untrained Qwen3.5-9B advisor, (b) the unconstrained trained advisor, and (c) the factuality-constrained trained advisor. Unconstrained training shifts strategy language toward technical and research registers, while factuality-constrained strategies use more news-compatible framing terms. Direct/rewriter-only models have no advisor strategies.

tion gains that transfer across model families, languages, and content domains. While unconstrained optimization frequently leads to unfaithful selection hacking, we further show that incorporating a factuality constraint helps preserve the original factual core while maintaining a substantial selection advantage. Overall, our findings show that language-model-mediated selection can be shaped by presentation alone, underscoring the need to evaluate AI agents across factuality, robustness, and user-aligned utility before deployment.

Limitations

First, real interfaces contain layout, ranking position, images, thumbnails, metadata, and interaction affordances that can affect both human and agent choices. We focus only on the textual presentation of an item, namely its title and abstract, and therefore do not model the full multimodal and interface-dependent nature of selection. Second, and more importantly, our experiments assume that the candidate list is already fixed before rewriting. This isolates a post-retrieval selection problem, but it does not test whether the rewritten item would still be retrieved or ranked into the same list by an upstream retrieval system. In a full pipeline, rewriting could change lexical or semantic matching and may therefore help or hurt retrieval under the original query or recommendation context. Our findings should therefore be interpreted as evidence of selection effects conditional on exposure, not as evidence of end-to-end optimization over retrieval, ranking, and selection. Third, although the impressions come from real logs, we treat each impression as a fixed snapshot. Real information-access systems are dynamic and non-stationary: articles enter and leave the corpus, competing stories evolve, user interests drift, social trends change, and platform feedback affects future retrieval and ranking. The selection pressures observed in a fixed candidate list may therefore change over time. Studying how rewriting incentives evolve under changing candidate pools and long-run feedback loops is important future work.

References

- Pranjal Aggarwal, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, Karthik Narasimhan, and Ameet Deshpande. 2024. [GEO: generative engine optimization](#). In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD 2024, Barcelona, Spain, August 25-29, 2024*, pages 5–16. ACM.
- Anthropic. 2025. Introducing Claude Haiku 4.5. <https://www.anthropic.com/news/claude-haiku-4-5>. Anthropic model announcement.
- Anthropic. 2026. Introducing Claude Sonnet 4.6. <https://www.anthropic.com/news/claude-sonnet-4-6>. Anthropic model announcement.
- Sagnik Anupam, Davis Brown, Shuo Li, Eric Wong, Hamed Hassani, and Osbert Bastani. 2025. [Browserarena: Evaluating llm agents on real-world web navigation tasks](#). *ArXiv preprint*, abs/2510.02418.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2024. [Self-rag: Learning to retrieve, generate, and critique through self-reflection](#). In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.
- Parth Asawa, Alan Zhu, Abby O’Neill, Matei Zaharia, Alexandros G. Dimakis, and Joseph E. Gonzalez. 2025. [How to train your advisor: Steering black-box llms with advisor models](#).
- Puneet S Bagga, Vivek F Farias, Tamar Korkotashvili, Tianyi Peng, and Yuhang Wu. 2025. [E-geo: A testbed for generative engine optimization in e-commerce](#). *ArXiv preprint*, abs/2511.20867.
- Bespoke Labs and Liyan Tang. 2024. Bespoke-MiniCheck-7B. <https://huggingface.co/BespokeLabs/Bespoke-MiniCheck-7B>. Hugging Face model card. Accessed: 2026-05-24.
- Arman Cohan, Sergey Feldman, Iz Beltagy, Doug Downey, and Daniel Weld. 2020. [SPECTER: Document-level representation learning using citation-informed transformers](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 2270–2282, Online. Association for Computational Linguistics.
- Paul Covington, Jay Adams, and Emre Sargin. 2016. [Deep neural networks for youtube recommendations](#). In *Proceedings of the 10th ACM Conference on Recommender Systems, Boston, MA, USA, September 15-19, 2016*, pages 191–198. ACM.
- Yixuan Du, Chenxiao Yu, Haoyan Xu, Ziyi Wang, Yue Zhao, and Xiyang Hu. 2026. [Multimodal generative engine optimization: Rank manipulation for vision-language model rankers](#). *ArXiv preprint*, abs/2601.12263.
- Lisa Dunlap, Trevor Darrell, Jacob Steinhardt, and Joseph E. Gonzalez. 2025. Stringsight: Turning model traces into actionable insights. <https://stringsight.com>. Accessed: 2025-01-26.
- Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. [Enabling large language models to generate text with citations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6465–6488, Singapore. Association for Computational Linguistics.
- Google. 2026. Gemini 3 Developer Guide. <https://ai.google.dev/gemini-api/docs/gemini-3>. Google AI for Developers documentation.
- Tyler Griggs, Sumanth Hegde, Eric Tang, Shu Liu, Shiyi Cao, Dacheng Li, Charlie Ruan, Philipp Moritz, Kourosh Hakhmaneshi, Richard Liaw, Akshay Malik, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. 2025. Evolving skylrl into a highly-modular rl framework. Notion Blog.

- Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. 2023. Atlas: Few-shot learning with retrieval augmented language models. *Journal of Machine Learning Research*, 24(251):1–43.
- Tanqiu Jiang, Min Bai, Nikolaos Pappas, Yanjun Qi, and Sandesh Swamy. 2025. [Cross-modal content optimization for steering web agent preferences](#). *ArXiv preprint*, abs/2510.03612.
- Haibo Jin, Ruoxi Chen, Peiyan Zhang, Yifeng Luo, Huimin Zeng, Man Luo, and Haohan Wang. 2026. [Controlling output rankings in generative engines for llm-based search](#). *ArXiv preprint*, abs/2602.03608.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. [Dense passage retrieval for open-domain question answering](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6769–6781, Online. Association for Computational Linguistics.
- Omar Khattab and Matei Zaharia. 2020. [Colbert: Efficient and effective passage search via contextualized late interaction over BERT](#). In *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval, SIGIR 2020, Virtual Event, China, July 25-30, 2020*, pages 39–48. ACM.
- Johannes Kruse, Kasper Lindschow, Saikishore Kalloori, Marco Polignano, Claudio Pomo, Abhishek Srivastava, Anshuk Uppal, Michael Riis Andersen, and Jes Frellsen. 2024. [EB-NeRD: A large-scale dataset for news recommendation](#). In *Proceedings of the Recommender Systems Challenge 2024*, pages 1–11. Association for Computing Machinery.
- Aounon Kumar and Himabindu Lakkaraju. 2024. [Manipulating large language models to increase product visibility](#).
- Komal Kumar, Aman Chadha, Salman Khan, Fahad Shahbaz Khan, and Hisham Cholakkal. 2026. [Paper circle: An open-source multi-agent research discovery and analysis framework](#). *ArXiv preprint*, abs/2604.06170.
- Patrick S. H. Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. [Retrieval-augmented generation for knowledge-intensive NLP tasks](#). In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*.
- Long Li, Weiwen Xu, Jiayan Guo, Ruochen Zhao, Xingxuan Li, Yuqian Yuan, Boqiang Zhang, Yuming Jiang, Yifei Xin, Ronghao Dang, and 1 others. 2024. [Chain of ideas: Revolutionizing research via novel idea development with llm agents](#). *ArXiv preprint*, abs/2410.13185.
- Xinyi Li, Yongfeng Zhang, and Edward C Malthouse. 2023. [A preliminary study of chatgpt on news recommendation: Personalization, provider fairness, fake news](#). *ArXiv preprint*, abs/2306.10702.
- Zeyi Liao, Lingbo Mo, Chejian Xu, Mintong Kang, Jiawei Zhang, Chaowei Xiao, Yuan Tian, Bo Li, and Huan Sun. 2025. Eia: Environmental injection attack on generalist web agents for privacy leakage. In *International Conference on Learning Representations*, volume 2025, pages 66972–67003.
- Dairui Liu, Boming Yang, Honghui Du, Derek Greene, Neil Hurley, Aonghus Lawlor, Ruihai Dong, and Irene Li. 2024. Recprompt: A self-tuning prompting framework for news recommendation using large language models. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 3902–3906.
- Nelson Liu, Tianyi Zhang, and Percy Liang. 2023. [Evaluating verifiability in generative search engines](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 7001–7025, Singapore. Association for Computational Linguistics.
- Abdelrahman Abdallah Bhawna Piryani Jamshid Mozafari and Mohammed Ali Adam Jatowt. 2025. How good are llm-based rerankers? an empirical analysis of state-of-the-art reranking models.
- Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, and 1 others. 2021. [Webgpt: Browser-assisted question-answering with human feedback](#). *ArXiv preprint*, abs/2112.09332.
- OpenAI. 2025. GPT-5 mini Model. <https://developer.openai.com/api/docs/models/gpt-5-mini>. OpenAI API model documentation.
- OpenAI. 2026. GPT-5.5 Model. <https://developer.openai.com/api/docs/models/gpt-5.5>. OpenAI API model documentation.
- Nilay Patel, Shivashankar Subramanian, Siddhant Garg, Pratyay Banerjee, and Amita Misra. 2024. [Towards improved multi-source attribution for long-form answer generation](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3906–3919, Mexico City, Mexico. Association for Computational Linguistics.
- Samuel Pfrommer, Yatong Bai, Tanmay Gautam, and Somayeh Sojoudi. 2024. [Ranking manipulation for conversational search engines](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 9523–9552, Miami, Florida, USA. Association for Computational Linguistics.

- Ronak Pradeep, Sahel Sharifmoghaddam, and Jimmy Lin. 2023a. [Rankvicuna: Zero-shot listwise document reranking with open-source large language models](#). *ArXiv preprint*, abs/2309.15088.
- Ronak Pradeep, Sahel Sharifmoghaddam, and Jimmy Lin. 2023b. [Rankzephyr: Effective and robust zero-shot listwise reranking is a breeze!](#) *ArXiv preprint*, abs/2312.02724.
- Jirui Qi, Gabriele Sarti, Raquel Fernández, and Arianna Bisazza. 2024. Model internals-based answer attribution for trustworthy retrieval-augmented generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 6037–6053.
- Zhen Qin, Rolf Jagerman, Kai Hui, Honglei Zhuang, Junru Wu, Le Yan, Jiaming Shen, Tianqi Liu, Jialu Liu, Donald Metzler, Xuanhui Wang, and Michael Bendersky. 2024. [Large language models are effective text rankers with pairwise ranking prompting](#). In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 1504–1518, Mexico City, Mexico. Association for Computational Linguistics.
- Qwen Team. 2026. [Qwen3.5: Towards native multi-modal agents](#).
- Claire E. Robertson, Nicolas Pröllochs, Kaoru Schwarzenegger, Philip Pärnamets, Jay J. Van Bavel, and Stefan Feuerriegel. 2023. [Negativity drives online news consumption](#). *Nature Human Behaviour*, 7:812–822.
- Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. 2024. [RAPTOR: recursive abstractive processing for tree-organized retrieval](#). In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. 2024. [Deepseekmath: Pushing the limits of mathematical reasoning in open language models](#). *ArXiv preprint*, abs/2402.03300.
- Sahel Sharifmoghaddam, Ronak Pradeep, Andre Slavesco, Ryan Nguyen, Andrew Xu, Zijian Chen, Yilin Zhang, Yidi Chen, Jasper Xian, and Jimmy Lin. 2025. Rankllm: A python package for reranking with llms. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3681–3690.
- Jiawen Shi, Zenghui Yuan, Guiyao Tie, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. 2025. [Prompt injection attack to tool selection in llm agents](#). *ArXiv preprint*, abs/2504.19793.
- Amanpreet Singh, Mike D’Arcy, Arman Cohan, Doug Downey, and Sergey Feldman. 2023. [SciRepEval: A multi-format benchmark for scientific document representations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5548–5566, Singapore. Association for Computational Linguistics.
- Weiwei Sun, Lingyong Yan, Xinyu Ma, Shuaiqiang Wang, Pengjie Ren, Zhumin Chen, Dawei Yin, and Zhaochun Ren. 2023. [Is ChatGPT good at search? investigating large language models as re-ranking agents](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 14918–14937, Singapore. Association for Computational Linguistics.
- Liyan Tang, Philippe Laban, and Greg Durrett. 2024a. [Minicheck: Efficient fact-checking of llms on grounding documents](#).
- Liyan Tang, Philippe Laban, and Greg Durrett. 2024b. MiniCheck-Flan-T5-Large. <https://huggingface.co/lytang/MiniCheck-Flan-T5-Large>. Hugging Face model card. Accessed: 2026-05-24.
- Yiming Tang, Yi Fan, Chenxiao Yu, Tiankai Yang, Yue Zhao, and Xiyang Hu. 2025. [Stealthrank: Llm ranking manipulation via stealthy prompt optimization](#). *ArXiv preprint*, abs/2504.05804.
- Zhihua Tian, Yuhao Chen, Yao Tang, Jian Liu, and Ruoxi Jia. 2026. [Diagnosing and repairing citation failures in generative engine optimization](#). *ArXiv preprint*, abs/2603.09296.
- Chong Wang and David M. Blei. 2011. [Collaborative topic modeling for recommending scientific articles](#). In *Proceedings of the 17th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Diego, CA, USA, August 21-24, 2011*, pages 448–456. ACM.
- Zihan Wang, Rui Zhang, Yu Liu, Wenshu Fan, Wenbo Jiang, Qingchuan Zhao, Hongwei Li, and Guowen Xu. 2026. Mpma: Preference manipulation attack against model context protocol. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 35838–35846.
- Chuhan Wu, Fangzhao Wu, Suyu Ge, Tao Qi, Yongfeng Huang, and Xing Xie. 2019. [Neural news recommendation with multi-head self-attention](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 6389–6394, Hong Kong, China. Association for Computational Linguistics.
- Fangzhao Wu, Ying Qiao, Jiun-Hung Chen, Chuhan Wu, Tao Qi, Jianxun Lian, Danyang Liu, Xing Xie, Jianfeng Gao, Winnie Wu, and Ming Zhou. 2020. [MIND: A large-scale dataset for news recommendation](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3597–3606, Online. Association for Computational Linguistics.

Yujia Wu, Shanshan Zhong, Yubin Kim, and Chenyan Xiong. 2025. [What generative search engines like and how to optimize web content cooperatively](#).

Sirui Xia, Xintao Wang, Jiaqing Liang, Yifei Zhang, Weikang Zhou, Jiaji Deng, Fei Yu, and Yanghua Xiao. 2025. Ground every sentence: Improving retrieval-augmented llms with interleaved reference-claim generation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 969–988.

Chejian Xu, Mintong Kang, Jiawei Zhang, Zeyi Liao, Lingbo Mo, Mengqi Yuan, Huan Sun, and Bo Li. 2024. [Advagent: Controllable blackbox red-teaming on web agents](#). *ArXiv preprint*, abs/2410.17401.

Yingxuan Yang, Mulei Ma, Yuxuan Huang, Huacan Chai, Chenyu Gong, Haoran Geng, Yuanjian Zhou, Ying Wen, Meng Fang, Muhao Chen, Shangding Gu, Ming Jin, Costas Spanos, Yang Yang, Pieter Abbeel, Dawn Song, Weinan Zhang, and Jun Wang. 2025. [Agentic web: Weaving the next web with ai agents](#). *Preprint*, arXiv:2507.21206.

Yue Yu, Wei Ping, Zihan Liu, Boxin Wang, Jiaxuan You, Chao Zhang, Mohammad Shoeybi, and Bryan Catanzaro. 2024. [Rankrag: Unifying context ranking with retrieval-augmented generation in llms](#). In *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*.

Weizhi Zhang, Yangning Li, Yuanchen Bei, Junyu Luo, Guancheng Wan, Liangwei Yang, Chenxuan Xie, Yuyao Yang, Wei-Chieh Huang, Chunyu Miao, and 1 others. 2025a. [From web search towards agentic deep research: Incentivizing search with reasoning agents](#). *ArXiv preprint*, abs/2506.18959.

Yanzhe Zhang, Tao Yu, and Diyi Yang. 2025b. Attacking vision-language computer agents via pop-ups. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8387–8401.

Zexuan Zhong, Ziqing Huang, Alexander Wettig, and Danqi Chen. 2023. [Poisoning retrieval corpora by injecting adversarial passages](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 13764–13775, Singapore. Association for Computational Linguistics.

Shengyao Zhuang, Honglei Zhuang, Bevan Koopman, and Guido Zuccon. 2024. [A setwise approach for effective and highly efficient zero-shot ranking with large language models](#). In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2024, Washington DC, USA, July 14-18, 2024*, pages 38–47. ACM.

Appendix

Our appendix is organized as follows:

- [Appendix A](#) discusses rewriting variants and the fixed-list intervention paradigm.
- [Appendix B](#) provides detailed data construction workflows and dataset split statistics for MIND, EB-NeRD, and SciRepEval.
- [Appendix C](#) presents the complete evaluation prompts, candidate rendering formats, and parsing rule specifications.
- [Appendix D](#) outlines reproducibility details, including hyperparameters, policy optimization configurations, and the core hardware/software environment.
- [Appendix E](#) evaluates verifier robustness alongside lexical and character-level overlap diagnostics across multiple target languages.
- [Appendix F](#) breaks down the qualitative strategy taxonomy and family distributions across optimized advisor variants.
- [Appendix G](#) details additional qualitative case studies covering unconstrained rewrites, factuality-aware constraints, and typical failure modes.
- [Appendix H](#) reports the machine translation protocols along with representative cross-lingual text evaluation examples.
- [Appendix I](#) documents data sources, data-use alignment, user privacy safeguards, and asset release policies.
- [Appendix J](#) states our AI Use Disclosure.

A Rewriting Variants and Fixed-List Intervention

This section provides diagrams for the rewriting setting compared in [section 3](#). All settings operate on the same fixed candidate list: only the target item may be rewritten, and all non-target items remain unchanged.

A.1 Original Text

The target agent receives the original candidate list without rewriting. This setting provides the reference target selection rate.

A.2 Rewriter-Only

In the standalone rewriter setting, the rewriter directly generates the rewritten title and abstract for the target item. We compare an untrained standalone rewriter, which uses a frozen rewriter model with a fixed prompt, and a trained standalone rewriter, which directly optimizes the rewriter using the target agent reward.

A.3 Untrained Advisor

An untrained advisor first produces a rewriting strategy for the target item. The same frozen rewriter then uses this strategy to generate the rewritten title and abstract. This baseline controls for the advisor-rewriter pipeline without selection-based learning.

A.4 Advisor-Guided Rewriter

In the advisor-guided setting, an advisor first produces a rewriting strategy for the target item, and a frozen rewriter then uses this strategy to generate the rewritten title and abstract.

We compare an untrained advisor, which controls for the advisor-rewriter pipeline without selection-based learning, and a trained advisor, which optimizes only the advisor using the target agent reward.

B Data Construction and Dataset Statistics

We use MINDlarge ([Wu et al., 2020](#)), EB-NeRD ([Kruse et al., 2024](#)), and SciRepEval-derived scientific-document candidate sets ([Singh et al., 2023](#)). We also construct translated evaluation-only variants of held-out MIND and EB-NeRD slates using Google Cloud Translation Advanced v3. This appendix summarizes the data sources, derived evaluation settings, intended use, and dataset statistics used in our experiments.

B.1 MIND

We use MINDlarge as the source dataset, but the original MINDlarge release split is not identical to the splits used in our experiments. We construct the primary news setting from MINDlarge impression logs. We keep impression-level candidate lists with at most 15 candidate items for computational tractability. Each constructed row contains one impression, one randomly sampled treated target from that impression, the target title and abstract, and the full candidate slate in the original order.

The advisor-training split uses 8,000 MIND impressions and the validation split uses 500 impressions. To reduce domination by highly repeated articles, the training construction caps treated-target reuse at two appearances per document. The final held-out MIND evaluation set uses 1,000 test impressions with a stricter cap of one treated occurrence per document, so no evaluation target repeats.

For leakage control, the held-out MIND evaluation set is separated from the training and validation sets by checking both document identifiers and rendered title–abstract text. The checked MIND paper-test split has zero overlap with the MIND train/validation splits for target IDs, target text, candidate IDs, and candidate text.

Table B.1: **MIND split statistics.** Rows denote impression-level candidate lists. Candidate-list sizes are computed from the fixed chooser candidate lists. The random baseline is the row-wise average probability of selecting the target item under a uniform random target agent, $\frac{1}{N} \sum_i 1/|L_i|$.

Split	Rows	Mean size	Max size	Random
Train	8,000	9.09	15	15.6%
Validation	500	9.36	15	15.0%
Test	1,000	8.42	15	16.9%

B.2 Translated MIND

We construct translated MIND variants as evaluation-only versions of the same 1,000 held-out MIND slates. These variants are used to test whether selection gains transfer across display languages while keeping the underlying candidate-slate structure fixed. Details of the translation procedure and preserved fields are provided in Section H.1.

B.3 EB-NeRD

We construct an EB-NeRD evaluation set from Danish news impressions. We sample 1,000 test impressions from records with at most 26 candidates and preserve full candidate lists rather than truncating them. Treated targets are sampled from each candidate slate with a cap of two treated occurrences per document. The resulting Danish EB-NeRD set has 886 unique treated documents, 114 repeated treated documents, and a maximum treated-target repeat count of two.

We also use an English machine-translated EB-NeRD variant. The Danish and English EB-NeRD versions preserve the same rows, impression IDs,

Table B.2: **Evaluation dataset statistics.** Rows are impression-level candidate lists. Mean and max report candidate-list size. Random is the row-wise uniform target-selection baseline, $\frac{1}{N} \sum_i 1/|L_i|$. For MIND, the English test set and its Arabic, Danish, Spanish, Swahili, Turkish, and Chinese translations use the same candidate lists. For EB-NeRD, the Danish set and its English machine-translated version use the same candidate lists.

Dataset	Rows	Mean	Max	Random
MIND test	1,000 each	8.42	15	16.9%
EB-NeRD	1,000 each	9.93	26	12.4%
SciRepEval	1,000	9.29	10	11.2%

treated article IDs, candidate slates, candidate order, and row alignment; only title and abstract text are translated.

B.4 SciRepEval-Derived Scientific-Document Candidate Lists

The scientific-document setting is derived from SciRepEval search-style candidate lists. Each example contains a fixed list of scientific-document candidates represented by titles and abstracts. The target document is selected from the candidate list and rewritten under the same fixed-slate intervention as the news settings.

C Prompts

Table C.1 reports the prompt templates used in our experiments. All generation prompts require valid JSON outputs to reduce parsing ambiguity. Empty title or abstract fields are rendered as None provided.

D Reproducibility Details

D.1 Evaluation Protocol

Our evaluation is conditional on exposure: for each held-out impression, the candidate list is fixed, and only the treated target item’s title and abstract are changed. Original and rewritten conditions use the same candidate order, same target agent prompt, and same target-selection metric.

D.2 Training Setup

Training uses GRPO through SkyRL with a Qwen3.5-9B policy. In the standalone setting, the policy directly generates the rewritten title and abstract. In the advisor-guided setting, the policy generates an advisor strategy that is passed to GPT-5-mini for the final rewrite. The reward is based on

Table C.1: **Prompt templates used in our experiments.** Panel A reports the advisor and rewriter prompts. Panel B reports the shared rewrite objective. Panel C reports the target-agent prompts. All candidate identifiers, candidate order, and non-target candidate texts are fixed across original and rewritten conditions.

Panel A: Advisor and rewriter prompts		
Component	System prompt	User prompt
Advisor	You are a rewriting strategist. Your job is to propose one concrete rewrite strategy. Do not write the final rewrite. Return valid JSON only.	Objective: {rewrite_objective} Current candidate title: {title} Current candidate abstract: {abstract} Return exactly one JSON object with this schema: { "strategy": "<one concrete open-ended instruction>". }
Advisor-guided rewriter	You are an execution model for rewriting. Your job is to execute the provided advisor strategy as directly as possible and rewrite the title and abstract. Return valid JSON only with rewritten_title and rewritten_abstract.	Objective: {rewrite_objective} Current candidate title: {title} Current candidate abstract: {abstract} Advisor strategy: {strategy} Return exactly one JSON object with this schema: { "rewritten_title": "...", "rewritten_abstract": "...". }
Standalone rewriter	You are a strategic text rewriter. Your job is to rewrite the title and abstract to maximize selection likelihood. Return valid JSON only with rewritten_title and rewritten_abstract.	Objective: {rewrite_objective} Current candidate title: {title} Current candidate abstract: {abstract} Return exactly one JSON object with this schema: { "rewritten_title": "...", "rewritten_abstract": "...". }

Panel B: Shared rewriter objective	
Objective	Prompt text
Rewrite objective	Rewrite the title and abstract to make this item more likely to be selected by the target agent based on its title and abstract. You may revise the title and abstract by changing their wording, emphasis, ordering, framing, clarity, structure, and tone when helpful.

Panel C: Target-agent prompts		
Setting	System prompt	User prompt
News	In this step, you are viewing a content feed. Each content card is clickable and has a unique label. You will be shown labeled items from the same impression. Each has: label, title, and abstract.	Items: {candidates} Choose exactly one item to click next, based only on the title and abstract. Return the selected label only.
Academic	In this step, you are viewing a set of candidate papers. Each paper is clickable and has a unique label. You will be shown labeled items from the same search session. Each has: label, title, and abstract.	Items: {candidates} Choose exactly one item to click next, based only on the title and abstract. Return the selected label only.

Panel D: Candidate rendering	
Field	Rendered format
Candidate list	A. Title: {title} Abstract: {abstract} B. Title: {title} Abstract: {abstract} ... Empty title or abstract fields are rendered as None provided.

Item	Setting
GPUs	4 NVIDIA A100-SXM4-80GB GPUs
Training framework	SkyRL
Core runtime	PyTorch and vLLM
Python	3.12.3
PyTorch	2.10.0
Transformers	5.6.0.dev0
vLLM	0.19.1
SkyRL	0.3.0
SkyRL Train	0.3.1
FlashAttention	2.8.3
Ray	2.51.1

Table D.1: Hardware and core software environment.

Table E.1: **BLEU overlap diagnostics on English inputs.** BLEU measures token n-gram overlap between the original and rewritten text. We report BLEU as a percentage and use it as a supplementary lexical-overlap diagnostic within each dataset rather than as a cross-language factuality measure.

Condition	Academic	MIND-En	EB-En
Untrained Rewriter-Only	15.0	15.7	11.9
Untrained Advisor	11.8	12.6	9.5
Rewriter-Only	19.0	1.2	1.2
Rewriter-Only + MiniCheck	21.2	15.0	12.4
Advisor	7.6	1.3	1.0
Advisor + MiniCheck	11.0	8.0	5.8

final target selection by the target agent. Source-support variants add a factuality-oriented support constraint.

D.3 Hardware and Software

Training was run with SkyRL, PyTorch, and vLLM on 4 NVIDIA A100-SXM4-80GB GPUs. The logged 9B unconstrained advisor run used 250 training batches.

E Verifier Robustness and Factuality Diagnostics

We use automatic factuality scores as diagnostics rather than definitive factuality labels. The factuality-constrained advisor was trained with MiniCheck-Flan-T5-Large (Tang et al., 2024a), so we use this verifier as the primary MiniCheck-based diagnostic for consistency with the training objective. To test whether the conclusions are specific to this verifier, we also evaluate Bespoke-MiniCheck-7B (Bespoke Labs and Tang, 2024) as a stronger independent grounded factuality checker. Table E.4 reports the Bespoke-MiniCheck-7B source-support scores on English MIND.

Table E.2: **chrF overlap diagnostics on English inputs.** chrF measures character n-gram overlap between the original and rewritten text. We use it as a lexical-overlap diagnostic rather than a factuality metric.

Condition	Academic	MIND-En	EB-En
Untrained Rewriter-Only	52.1	49.7	45.3
Untrained Advisor	48.9	45.4	41.7
Rewriter-Only	54.5	17.5	16.0
Rewriter-Only + MiniCheck	55.0	46.0	43.4
Advisor	44.5	18.8	16.1
Advisor + MiniCheck	49.1	39.6	35.3

The two verifiers produce similar aggregate trends, with differences of roughly two to three points. The direction of the difference varies by dataset: Bespoke-MiniCheck-7B is slightly more conservative on the news-style datasets, while assigning slightly higher scores on the academic-domain set. This pattern suggests verifier-specific calibration differences rather than a systematic reversal of the factuality diagnostics.

For this reason, we do not interpret either MiniCheck variant as a calibrated measure of factuality across all settings. Instead, we use MiniCheck scores to check whether the advisor’s gains are accompanied by large changes under grounded factuality verifiers. In the main experiments, the selection results remain the primary outcome, while MiniCheck provides an auxiliary check on factual-consistency risk.

We also report lexical-overlap diagnostics between the original and rewritten texts. Table E.1 reports BLEU on English inputs, and Table E.2 reports chrF on the same English settings. These metrics measure token- and character-level overlap, respectively, and should be interpreted as rewrite-distance diagnostics rather than factuality metrics.

For multilingual settings, we treat raw MiniCheck scores with additional caution. MiniCheck is primarily an English-oriented grounded factuality verifier, and raw scores can be poorly calibrated on non-English text. We therefore report chrF overlap across the MIND language variants in Table E.3. These lexical metrics characterize how much the rewrite differs from the original, but they do not establish factual consistency.

Table E.3: **chrF overlap diagnostics across MIND languages.** chrF measures character n-gram overlap between the original and rewritten text. We report chrF as a percentage and use it as a lexical-overlap diagnostic rather than a factuality metric.

Dataset	Untrained Rewriter-Only	Untrained Advisor	Rewriter-Only	Rewriter-Only + MC	Advisor	Advisor + MC
MIND-En	49.7	45.4	17.5	46.0	18.8	39.6
Arabic	49.9	44.3	0.2	49.5	18.5	40.0
Danish	53.1	47.7	15.4	53.5	20.1	42.8
Spanish	52.6	49.7	17.4	52.8	22.0	46.0
Swahili	58.5	48.2	12.9	58.2	22.0	41.9
Turkish	53.0	47.9	13.0	53.5	19.9	42.0
Chinese	31.4	27.8	0.3	33.9	4.6	24.0

Table E.4: **Bespoke-MiniCheck-7B source-support scores on English MIND.** Scores are aggregated over rewritten title and abstract claims for the English evaluation split ($N = 1,000$ per setting). Higher values indicate greater support by the source document.

Setting	Doc mean (std)
Standalone	0.003 (0.005)
Standalone + MiniCheck	0.386 (0.231)
Advisor	0.009 (0.026)
Advisor + MiniCheck	0.268 (0.144)
Base9B Direct	0.517 (0.183)
Base9B Advisor	0.357 (0.192)

F Strategy Taxonomy and Qualitative Examples

F.1 Strategy analysis with StringSight.

To characterize the rewriting behavior learned by the advisor, we conduct a post-hoc qualitative analysis using StringSight. We import per-example advisor rewrite outputs and associated strategy descriptions into StringSight, which provides an initial clustering of recurring rewriting behaviors. We treat these clusters as an exploratory taxonomy rather than as automatic ground-truth labels. We then merge and normalize the raw clusters into paper-facing strategy families, focusing on behavioral rewriting moves such as article-specific reframing, urgency or engagement framing, actionability, grounding, credibility, and factuality risk. All quantitative task results are computed independently of StringSight; StringSight is used only to support interpretation of the advisor’s learned strategies.

F.2 Advisor Strategy Taxonomy

Table F.1 summarizes the qualitative strategy taxonomy used to analyze advisor outputs. We code the advisor instruction rather than the final rewritten title or abstract, since the instruction is the object

directly optimized by the advisor policy. The taxonomy is derived from rule-normalized StringSight clusters and merged into paper-facing families that capture the main behavioral rewriting move, such as topic substitution, grounded expansion, stakes framing, evidential framing, or genre reframing. These categories describe mechanisms by which advisor instructions may increase target selection rate; they should not be interpreted as judgments that the resulting rewrites are more accurate, useful, or socially desirable.

All counts are computed on the English MIND evaluation set with one advisor instruction per example ($N = 1,000$ per variant). The no-support-constraint advisor is dominated by unsupported technical substitution, while the support-constrained advisor distributes more broadly across grounded expansion, urgency and stakes framing, mixed investigative reframing, and academic or empirical genre reframing. Evaluator-specific advisors also show distinct policies: the Haiku 4.5 evaluator advisor mostly expands article-grounded details, whereas the Gemini 3 Flash evaluator advisor mostly academicizes items into empirical-study formats.

G Additional Qualitative Examples

G.1 Qualitative Examples

This appendix provides additional qualitative examples from the English MIND evaluation split. Each example keeps the original candidate list fixed and rewrites only the target item. To keep the examples focused on the intervention, we omit competing item text and report only the originally selected non-target identifier and the chooser outcome before and after rewriting. We avoid political, discriminatory, sexual, and high-risk news examples in this section.

The examples illustrate three recurring behav-

Table F.1: **Qualitative taxonomy of advisor strategies.** Each row summarizes one paper-facing strategy family, its main rewriting move, and a representative advisor instruction.

Strategy family	Main rewriting move	Representative advisor instruction
Unsupported technical substitution	Replaces the source article’s actual topic with a new technical, scientific, or AI-centered problem. This can be attractive to the chooser but often introduces unsupported mechanisms or claims.	Reframe the content from a general economic warning into a specific technical proposal for a novel AI-driven financial forecasting model that solves the “black box” problem in recession prediction, including architecture and accuracy claims.
Academic or empirical genre re-framing	Recasts a news item as an academic paper, empirical study, policy analysis, or methodological contribution. The distinctive move is academicization through metrics, methods, data sources, or results.	Transform the title and abstract into a rigorous empirical study by defining a quantifiable metric for policy rigidity or investment inertia, adding statistical evidence, and presenting the rewrite in an IMRD-style structure.
Grounded article expansion and detailing	Keeps the core article topic fixed while making source-relevant details more explicit. Often expands a short abstract into a fuller explanation or foregrounds concrete actors and consequences.	Expand the abstract into a multi-paragraph structure that details policy failures, gives concrete examples of investment misalignment, and contrasts proposed changes with the current cyclical approach.
Urgency and stakes framing	Heightens perceived importance by foregrounding risk, rarity, human cost, consequences, or time sensitivity. The factual topic is usually preserved, but salience is shifted.	Reframe the narrative to emphasize the rarity and potential systemic safety concerns of a repeat crash at the same location, while highlighting the human cost of three injuries.
Authority and evidential framing	Makes the item feel more credible, investigative, evidence-oriented, neutral, or institutionally grounded. Often invokes legal, forensic, policy, data-backed, or expert framing.	Reframe the content from a celebrity gossip update into an investigative feature story that highlights the legal and forensic complexities of the stolen diamond ring case.
Human and emotional impact framing	Foregrounds personal vulnerability, community impact, grief, resilience, identity, or lived experience. This makes a routine or abstract event feel personally consequential.	Reframe the narrative to emphasize the vulnerability of the Hmong community and the brazen nature of the attack, while replacing sensationalized language with more precise community-safety framing.
Positive benefit and agency re-framing	Makes the item more constructive by emphasizing resolution, safety, usefulness, agency, or practical benefit. Often reduces alarm while preserving selection appeal.	Reframe the narrative to emphasize the routine nature of the mechanical issue and the successful resolution, while stating that the former Secretary of State remained safe and unharmed.
Causal and structural clarity	Clarifies mechanisms, causes, systems, or structural implications behind an event. The emphasis is explanatory rather than merely evidential.	Reframe the report into an investigative feature that explores the medical mechanisms of the brain injury and broader safety implications for amateur combat sports in the UK.
Other or mixed strategy	Captures instructions that combine multiple moves or do not cleanly fit one primary category. Many examples mix investigative framing, historical significance, irony, and stakes.	Reframe the narrative into an investigative feature story that emphasizes the historical significance of the third consecutive gubernatorial rejection and the irony of parole being blocked after five decades.

iors. First, unconstrained optimization can achieve selection flips by changing the topic or genre of the target item. Second, MiniCheck-aware optimization usually preserves more source content while still improving target selection. Third, faithful rewrites are not always sufficient to overturn the chooser’s original preference under a fixed candidate list.

G.2 Case-Oriented English Rewriting Examples

This appendix gives qualitative examples of fixed-list rewriting behavior. Each example keeps the candidate list, candidate order, and all non-target items unchanged; only the target item’s title and abstract are rewritten. The examples are selected to illustrate recurring behaviors rather than to estimate their frequency.

G.3 Same-Target Comparison Across English Rewriter Variants

This appendix holds one target item and its candidate list fixed while comparing how six English rewriter variants rewrite the same source. The example shows that the variants do not merely produce stylistic paraphrases: some rewrites preserve the original visual-science story, while others shift toward scientific or technical framings with different source-support and selection outcomes.

H Translation Protocol and Examples

H.1 Translation Protocol

For MIND, we export unique document title–abstract pairs from the 1,000 held-out slates and translate them with Google Cloud Translation API (Advanced v3) using the default neural machine translation model. The target languages are Arabic,

Table F.2: **Post-hoc distribution of advisor strategy families.** Each column reports the number and percentage of advisor instructions assigned to each strategy family on the English MIND evaluation set ($N = 1,000$ per column). Strategy families are assigned automatically from the generated advisor instruction. The distributions show how factual-support constraints and transfer target agents change the dominant rewriting strategy.

Strategy family	Base	No supp.	Supp.-constr.	Haiku 4.5	Gemini 3 Flash
Unfaithful technical topic pivot	6 (0.6%)	962 (96.2%)	1 (0.1%)	0 (0.0%)	31 (3.1%)
Academic investigative genre reframing	28 (2.8%)	13 (1.3%)	67 (6.7%)	72 (7.2%)	951 (95.1%)
Grounded article expansion and detailing	553 (55.3%)	1 (0.1%)	456 (45.6%)	928 (92.8%)	4 (0.4%)
Urgency and stakes engagement framing	300 (30.0%)	13 (1.3%)	283 (28.3%)	0 (0.0%)	6 (0.6%)
Authority, evidence, and neutrality framing	44 (4.4%)	1 (0.1%)	41 (4.1%)	0 (0.0%)	1 (0.1%)
Human and emotional impact framing	18 (1.8%)	1 (0.1%)	27 (2.7%)	0 (0.0%)	0 (0.0%)
Positive benefit and agency reframing	21 (2.1%)	5 (0.5%)	8 (0.8%)	0 (0.0%)	2 (0.2%)
Causal and structural clarity	5 (0.5%)	2 (0.2%)	9 (0.9%)	0 (0.0%)	3 (0.3%)
Other or mixed strategy	25 (2.5%)	2 (0.2%)	108 (10.8%)	0 (0.0%)	2 (0.2%)
Total	1000 (100.0%)	1000 (100.0%)	1000 (100.0%)	1000 (100.0%)	1000 (100.0%)

Danish, Spanish, Swahili, Turkish, and Simplified Chinese. For EB-NeRD, we translate Danish title–abstract pairs into English.

Across translation-based evaluations, only the displayed candidate title and abstract are translated. Row order, impression IDs, treated candidate IDs, candidate counts, candidate order, and nested candidate-ID sequences are preserved. Candidate identifiers, labels, and prompt structure are not translated. Empty title or abstract fields are rendered as None provided. No additional training, fine-tuning, or translation-specific rewriting is performed.

This protocol tests language transfer under structurally aligned candidate slates. It does not claim to reproduce native language-specific presentation conventions, such as natural headline style, emphasis, or framing.

H.2 Representative Translation Examples

Tables H.1 and H.2 show two representative MIND target items after machine translation. The first is a longer informational item with numeric content, and the second is a shorter food and culture item. Across all examples, the candidate list, target item, and prompt structure are preserved; only the displayed title and abstract fields are translated.

Table G.1: **Qualitative fixed-list rewriting cases.** Each case compares the original target title with the rewritten target title under the same fixed candidate list. The examples illustrate source-supported gains, unsupported gains, and faithful rewrites that remain unselected. MiniCheck reports the source-support score for the rewritten target.

Case	Original target title	Rewritten target title	Model	Outcome
Grounded selection gain	Cold, wet weather rolls into western Washington, could lower snow levels to 1,000 feet by Tuesday	Holiday Travel Alert: Snow Levels Could Drop to 1,000 Feet; Expect Freezing, Hazardous Mountain Driving This Week	Before-training advisor	$0 \rightarrow 1$ MC: 0.868
<i>Interpretation.</i> The rewrite foregrounds travel risk while preserving the reported weather forecast.				
Unsupported topic pivot	Brave sea turtle shows up to visit scuba divers at shark dive	DeepSeaDetect: A Machine Learning Pipeline for Reliable Hawksbill Turtle Detection in Shark-Associated Midwater Environments	Trained advisor	$0 \rightarrow 1$ MC: 0.014
<i>Interpretation.</i> The unconstrained advisor wins selection by reframing a nature-video item as an unsupported technical paper.				
Source-supported selection gain	Alaskan city sees heat and snowfall records in single day	A Rare Weather Paradox: Anchorage Records a 45 degrees F High and Over a Foot of Snow in One Day	Trained advisor + MiniCheck	$0 \rightarrow 1$ MC: 0.643
<i>Interpretation.</i> The MiniCheck-aware advisor sharpens the weather anomaly while remaining close to the source.				
Faithful failure	Duke climbs to No. 1 in AP Top 25 following Kentucky’s loss	Duke Surges to No. 1 in AP Top 25 Poll Following Kentucky’s Loss	Trained standalone + MiniCheck	$0 \rightarrow 0$ MC: 0.960
<i>Interpretation.</i> A highly source-supported rewrite can still fail, showing that source support alone does not guarantee selection.				

Table G.2: **Same-target English comparison across rewriter variants.** The target item and candidate list are fixed; only the target title and abstract are rewritten. Selection denotes whether the rewritten target is selected by the target agent.

Variant	Target title	Selection	Support	Interpretation
Original target	Mesmerising moment bubble freezes instantly in sub-zero temperatures	0	–	The original target is a visual-interest news item and is not selected.
Before-training direct	Incredible Sub-Zero Phenomenon: Man Creates Instant Snowflake-Patterned Ice Bubble	0	MiniCheck: 0.862 chrF: 0.483	A close, source-supported rewrite improves presentation but does not change selection.
Before-training advisor	Instantaneous Crystallization: Bubble Freezes into Delicate Snowflake Patterns in Sub-Zero Nebraska Test	0	MiniCheck: 0.159 chrF: 0.399	The rewrite adopts a more technical frame but still remains unselected.
Trained standalone	Real-Time Visual Analysis of Rapid Phase Transitions in Supercooled Fluids	1	MiniCheck: 0.010 chrF: 0.323	Selection gain comes with a strong scientific reframing that is poorly supported by the source.
Trained advisor	Controlled Bubble-Freezing Technique Produces Stable Snowflake-Like Ice Structures in Sub-Zero Conditions	0	MiniCheck: 0.131 chrF: 0.332	The advisor shifts toward a method-like framing, but the target remains unselected.
Trained standalone + MiniCheck	Capturing Nature’s Art: Instant Snowflake Patterns Form on a Bubble Frozen in Sub-Zero Temperatures	1	MiniCheck: 0.033 chrF: 0.429	The rewrite is closer in wording than the unconstrained trained standalone case, but MiniCheck still assigns low support.
Trained advisor + MiniCheck	Instantly Frozen Bubble Reveals Rapid Crystallization and Intricate Snowflake Patterns	0	MiniCheck: 0.170 chrF: 0.391	The rewrite preserves the main event while using a sharper technical frame, but this does not guarantee selection.

Original target metadata: target N118672; impression 559; originally unselected. The original abstract describes a bubble freezing instantly with snowflake patterns in Nebraska on November 12; the table truncates abstracts to focus on title-level framing.

Table H.1: **Representative MIND multilingual translation example 1: airport punctuality.** Target candidate ID: N84539; fixed candidate list size: 9.

Language	Target item displayed to the chooser
English source	Title: Why Tokyo's Haneda is one of the world's most punctual airports Abstract: Haneda, officially called Tokyo International Airport, is the world's fifth busiest airport. Yet it delivered 85.6% of its flights on time in 2018, making it the most punctual mega airport in the world.
Arabic	Title: لماذا يُعد مطار هانيدا في طوكيو أحد أكثر المطارات التزاماً بالمواعيد في العالم؟ Abstract: يُعد مطار هانيدا، المعروف رسميًا باسم مطار طوكيو الدولي، خامس أكثر مطارات العالم ازدحامًا. ومع ذلك، فقد أنجز 85,6 بالمئة من رحلاته في الموعد المحدد عام 2018، مما يجعله أكثر المطارات الضخمة التزامًا بالمواعيد في العالم.
Danish	Title: Hvorfor Tokyos Haneda er en af verdens mest punktlige lufthavn Abstract: Haneda, officielt kaldet Tokyo International Airport, er verdens femte travleste lufthavn. Alligevel afviklede den 85,6% af sine flyvninger til tiden i 2018, hvilket gør den til den mest punktlige megalufthavn i verden.
Spanish	Title: ¿Por qué el aeropuerto de Haneda, en Tokio, es uno de los aeropuertos más puntuales del mundo? Abstract: Haneda, cuyo nombre oficial es Aeropuerto Internacional de Tokio, es el quinto aeropuerto con mayor tráfico del mundo. Sin embargo, en 2018 logró que el 85,6% de sus vuelos llegaran a tiempo, lo que lo convierte en el megaaeropuerto más puntual del mundo.
Swahili	Title: Kwa nini Haneda ya Tokyo ni mojawapo ya viwanja vya ndege vinavyofanya kazi kwa wakati zaidi duniani Abstract: Haneda, ambayo inaitwa rasmi Uwanja wa Ndege wa Kimataifa wa Tokyo, ni uwanja wa ndege wa tano wenye shughuli nyingi zaidi duniani. Hata hivyo, ilitoa 85.6% ya safari zake za ndege kwa wakati mwaka wa 2018, na kuifanya kuwa uwanja mkubwa wa ndege unaofanya kazi kwa wakati zaidi duniani.
Turkish	Title: Tokyo'nun Haneda Havalimanı neden dünyanın en dakik havalimanlarından biri? Abstract: Resmi adı Tokyo Uluslararası Havalimanı olan Haneda, dünyanın en yoğun beşinci havalimanıdır. Buna rağmen, 2018 yılında uçuşlarının %85,6'sını zamanında gerçekleştirerek dünyanın en dakik mega havalimanı olmuştur.
Chinese (zh-CN)	Title: 为什么东京羽田机场是世界上最准时的机场之一？ Abstract: 羽田机场，正式名称为东京国际机场，是世界第五繁忙的机场。然而，2018年其航班准点率高达85.6%，使其成为世界上最准点的巨型机场。

Table H.2: **Representative MIND multilingual translation example 2: food and culture.** Target candidate ID: N66002; fixed candidate list size: 2.

Language	Target item displayed to the chooser
English source	Title: 5 Things You Didn't Know About Campbell's Iconic Green Bean Casserole Abstract: We even got our hands on the original 1955 recipe card from the test kitchen.
Arabic	Title: خمسة أشياء لم تكن تعرفها عن طبق الفاصوليا الخضراء الشهير من كامبل Abstract: بل إننا حصلنا على بطاقة الوصفة الأصلية لعام 1955 من مطبخ الاختبار.
Danish	Title: 5 ting du ikke vidste om Campbells ikoniske grønne bønnegryde Abstract: Vi fik endda fat i det originale opskriftskort fra 1955 fra testkøkkenet.
Spanish	Title: 5 cosas que no sabías sobre la icónica cazuela de judías verdes de Campbell's Abstract: Incluso conseguimos la tarjeta de recetas original de 1955 de la cocina de pruebas.
Swahili	Title: Mambo 5 Uliyokuwa Huyajui Kuhusu Kitoweo cha Maharagwe Kijani cha Campbell Abstract: Tulipata hata kadi ya mapishi ya awali ya 1955 kutoka jikoni ya majaribio.
Turkish	Title: Campbell's'ın İkonik Yeşil Fasulye Güveci Hakkında Bilmediğiniz 5 Şey Abstract: Hatta test mutfağından 1955 yılına ait orijinal tarif kartını bile ele geçirdik.
Chinese (zh-CN)	Title: 关于坎贝尔经典青豆砂锅，你不知道的5件事 Abstract: 我们甚至还找到了1955年测试厨房的原始食谱卡。

I Data Sources, Licenses, and Release Policy

We use MINDlarge (Wu et al., 2020), EB-NeRD (Kruse et al., 2024), and SciRepEval-derived scientific-document candidate sets (Singh et al., 2023). We also construct translated evaluation-only variants of held-out MIND and EB-NeRD slates using Google Cloud Translation API.

Data sources and derived selection format. MINDlarge provides English news recommendation impressions and is used for advisor training and in-domain evaluation. EB-NeRD provides Danish news recommendation impressions from Ekstra Bladet and is used only for out-of-distribution news evaluation. For scientific-document evaluation, we use candidate-paper sets derived from the SciRepEval search setting. Across all settings, each example is converted into a fixed candidate-list selection instance: one target item may have its displayed title and abstract rewritten, while all other candidates remain unchanged. Translated MIND and EB-NeRD variants preserve the original candidate lists and target assignments, but replace the displayed candidate text with machine-translated text. These translated vari-

ants are used only for controlled evaluation.

Intended use. We use these datasets for controlled research on presentation-induced selection shifts under fixed exposure. Our experiments do not train or deploy a production recommendation system, do not model long-term platform feedback, and do not attempt to infer private user attributes. Although MIND and EB-NeRD contain interaction logs, we use them only to construct candidate-list selection contexts. The advisor is optimized using chooser-model feedback, not by predicting individual users’ clicks.

Licenses, access conditions, and release policy. We use all datasets and derived evaluation sets under their respective license and access terms. MINDlarge is made available for research purposes under the Microsoft Research License Terms. For MIND-derived settings, we release code, preprocessing scripts, prompts, target assignments, model outputs, and aggregate evaluation results, but do not redistribute raw MIND impression logs or article metadata; users must obtain MIND from the original provider and run our preprocessing scripts.

EB-NeRD is made available under the dataset’s General License Terms for research-only, non-commercial use. Because the EB-NeRD terms restrict redistribution of exact dataset contents and other derivative uses, including training language models, we use EB-NeRD only for controlled research evaluation. We do not train or fine-tune language-model, advisor, or rewriter parameters on EB-NeRD text, logs, or article metadata, and we do not release EB-NeRD text, logs, article metadata, translated EB-NeRD slates, or reconstructed EB-NeRD candidate-list corpora.

For the scientific-document experiments, we use candidate sets derived from SciRepEval. The SciRepEval aggregate benchmark is released under the Open Data Commons Attribution License (ODC-BY), and we follow the applicable licensing requirements of its constituent datasets. Where permitted by the SciRepEval license and constituent dataset licenses, we release the processed scientific-document evaluation files used in our experiments, with attribution to the original sources.

Translations preserve the access restrictions of the underlying source datasets. Accordingly, we release translated or processed artifacts only where permitted by the source dataset terms.

Privacy and content considerations. We do not collect new user data or conduct new human-subject data collection. MIND and EB-NeRD are existing datasets with anonymized user identifiers. We do not attempt to identify users, recover private attributes, link anonymized identifiers to external accounts, or reconstruct user-level histories beyond the candidate-list context needed for our experiments. When reporting qualitative examples, we avoid releasing raw user logs and exact EB-NeRD dataset contents.

Potential misuse. Our method studies how rewriting displayed text can increase selection by an LLM chooser, which could be misused to make content more appealing to automated selection systems without improving its underlying quality. We therefore frame the work as a controlled measurement of selection vulnerability, report factuality variants, and discuss limitations and risks in the limitations section.

J AI Use Disclosure

The authors used AI-based tools to assist with code generation, editing, and writing during the preparation of this paper. Specifically, AI assistance was used to help draft and revise portions of the manuscript for clarity, grammar, and organization, and to support the development, debugging, and refinement of code used in the research workflow. All AI-generated or AI-assisted content, code, analyses, and interpretations were reviewed, verified, and, where necessary, modified by the authors. The authors take full responsibility for the accuracy, integrity, originality, and final content of the paper, including any code or text developed with AI assistance.